

Sri Ram Sharma

sriramxdev@gmail.com | linkedin.com/in/sriramxdev | github.com/sriramxdev | sriramxdev.me

EDUCATION

United College of Engineering & Research
B.Tech, CSE (Artificial Intelligence & Machine Learning)

Prayagraj, UP, India
Expected 07/2027

St. Mary's Convent School
Senior Secondary (CBSE)

Prayagraj, UP, India
2019 – 2021

PROJECTS

Echo-Sync | *Python, FastAPI, PyTorch, ONNX, Docker/Podman* | [GitHub](#) July 2026 – Present

- Directing system architecture as technical lead for a 4-person team, defining REST API contracts between the inference service and client layer and standardizing Docker/Podman images to eliminate cross-machine build drift.
- Engineered a kinematic frame-pruning stage in Python that filters redundant MediaPipe landmark frames via velocity thresholding, cutting edge CPU load by 30%.
- Deployed an ONNX-quantized ST-GCN model for sub-100ms local inference, with Prometheus aggregating anonymized client metrics for monitoring.

DynaPrune | *Python, PyTorch, ONNX, Latency Profiling* | [GitHub](#) Sep 2026 – Present

- Developing a structured model-pruning pipeline in PyTorch to compress the ST-GCN model used in Echo-Sync, targeting parameter-count reduction without exceeding a defined accuracy-drop threshold.
- Benchmarking channel-level sparsity ratios against inference latency, exporting pruned checkpoints to ONNX to measure real-world speedup over the unpruned baseline.
- Building a reproducible profiling harness to compare pre- and post-pruning latency and memory footprint across CPU and edge-constrained environments.

PDF Document Parser | *Python, PyMuPDF, Streamlit, JSON* | [GitHub](#) Nov 2025 – Jan 2026

- Engineered a document extraction pipeline using PyMuPDF to extract document-level metadata, page-wise token/character metrics, and raw text into standardized JSON schemas.
- Implemented multi-layer input validation, including a magic-byte header check (%PDF-) and a 25 MB memory cap to block spoofed, encrypted, or corrupted files.
- Built an interactive, dual-theme Streamlit interface supporting interactive JSON tree previews and client-side exports, deployed on Streamlit Cloud.

A-Stride | *Kotlin, Android SDK (API 29+), MVVM* | [GitHub](#) March 2026 – Present

- Architected a privacy-focused Android step-tracking engine running entirely on-device, with zero network permissions requested.
- Interfaced directly with hardware step-counter sensors (TYPE_STEP_COUNTER) under an MVVM architecture, avoiding polling overhead that drains battery.
- Calibrated a 0.73m stride matrix via controlled-course field testing, achieving 98.4% distance-estimation precision.

TECHNICAL SKILLS

Languages: Python, Java, Rust, C/C++, SQL, HTML/CSS

Frameworks/Tools: PyTorch, MediaPipe, OpenCV, Scikit-Learn, NumPy, Pandas, FastAPI, LangChain, LlamaIndex, Hugging Face

Databases & Data Systems: PostgreSQL, MongoDB, MySQL, Firebase

Platforms/DevOps: Linux (Fedora/Ubuntu), Docker, Podman, Git, Prometheus, Grafana

CERTIFICATIONS

GitHub Foundations Certification | [Microsoft/GitHub](#)
CRUD Operations in MongoDB | [Credly Verified](#)

Aug 2026 – Aug 2028
Aug 2026